

2020320064 문정민(컴퓨터학과)

7.1. From Fully Connected Layers to Convolutions

no code

7.2. Convolutions for Images


!pip install d2l==1.0.3

 숨겨진 출력 표시

```
import torch
from torch import nn
from d2l import torch as d2l

def corr2d(X, K):
    """Compute 2D cross-correlation."""
    h, w = K.shape
    Y = torch.zeros((X.shape[0] - h + 1, X.shape[1] - w + 1))
    for i in range(Y.shape[0]):
        for j in range(Y.shape[1]):
            Y[i, j] = (X[i:i + h, j:j + w] * K).sum()
    return Y


X = torch.tensor([[0.0, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]])
K = torch.tensor([[0.0, 1.0], [2.0, 3.0]])
corr2d(X, K)
```

 tensor([[19., 25.],
 [37., 43.]])

```
class Conv2D(nn.Module):
    def __init__(self, kernel_size):
        super().__init__()
        self.weight = nn.Parameter(torch.rand(kernel_size))
        self.bias = nn.Parameter(torch.zeros(1))


    def forward(self, x):
        return corr2d(x, self.weight) + self.bias
```

```
X = torch.ones((6, 8))
X[:, 2:6] = 0
X
```


 tensor([[1., 1., 0., 0., 0., 0., 1., 1.],
 [1., 1., 0., 0., 0., 0., 1., 1.],
 [1., 1., 0., 0., 0., 0., 1., 1.],
 [1., 1., 0., 0., 0., 0., 1., 1.],
 [1., 1., 0., 0., 0., 0., 1., 1.],
 [1., 1., 0., 0., 0., 0., 1., 1.]])

```
K = torch.tensor([[1.0, -1.0]])
```

```
Y = corr2d(X, K)
Y
```

 tensor([[0., 1., 0., 0., 0., -1., 0.],
 [0., 1., 0., 0., 0., -1., 0.],
 [0., 1., 0., 0., 0., -1., 0.],
 [0., 1., 0., 0., 0., -1., 0.],
 [0., 1., 0., 0., 0., -1., 0.],
 [0., 1., 0., 0., 0., -1., 0.]])

```
corr2d(X.t(), K)
```

 tensor([[0., 0., 0., 0., 0.],
 [0., 0., 0., 0., 0.],
 [0., 0., 0., 0., 0.],
 [0., 0., 0., 0., 0.]])

```
[0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0.]])
```

```
# Construct a two-dimensional convolutional layer with 1 output channel and a
# kernel of shape (1, 2). For the sake of simplicity, we ignore the bias here
conv2d = nn.LazyConv2d(1, kernel_size=(1, 2), bias=False)
```

```
# The two-dimensional convolutional layer uses four-dimensional input and
# output in the format of (example, channel, height, width), where the batch
# size (number of examples in the batch) and the number of channels are both 1
X = X.reshape((1, 1, 6, 8))
Y = Y.reshape((1, 1, 6, 7))
lr = 3e-2 # Learning rate
```

```
for i in range(10):
    Y_hat = conv2d(X)
    l = (Y_hat - Y) ** 2
    conv2d.zero_grad()
    l.sum().backward()
    # Update the kernel
    conv2d.weight.data[:] -= lr * conv2d.weight.grad
    if (i + 1) % 2 == 0:
        print(f'epoch {i + 1}, loss {l.sum():.3f}')
```

```
epoch 2, loss 10.132
epoch 4, loss 2.001
epoch 6, loss 0.459
epoch 8, loss 0.128
epoch 10, loss 0.042
```

```
conv2d.weight.data.reshape((1, 2))
```

```
tensor([[ 0.9652, -1.0038]])
```

✓ 7.3. Padding and Stride

```
import torch
from torch import nn
```

```
# We define a helper function to calculate convolutions. It initializes the
# convolutional layer weights and performs corresponding dimensionality
# elevations and reductions on the input and output
def comp_conv2d(conv2d, X):
    # (1, 1) indicates that batch size and the number of channels are both 1
    X = X.reshape((1, 1) + X.shape)
    Y = conv2d(X)
    # Strip the first two dimensions: examples and channels
    return Y.reshape(Y.shape[2:])
```

```
# 1 row and column is padded on either side, so a total of 2 rows or columns
# are added
conv2d = nn.LazyConv2d(1, kernel_size=3, padding=1)
X = torch.rand(size=(8, 8))
comp_conv2d(conv2d, X).shape
```

```
torch.Size([8, 8])
```

```
# We use a convolution kernel with height 5 and width 3. The padding on either
# side of the height and width are 2 and 1, respectively
conv2d = nn.LazyConv2d(1, kernel_size=(5, 3), padding=(2, 1))
comp_conv2d(conv2d, X).shape
```

```
torch.Size([8, 8])
```

```
conv2d = nn.LazyConv2d(1, kernel_size=3, padding=1, stride=2)
comp_conv2d(conv2d, X).shape
```

```
torch.Size([4, 4])
```

```
conv2d = nn.LazyConv2d(1, kernel_size=(3, 5), padding=(0, 1), stride=(3, 4))
comp_conv2d(conv2d, X).shape
```

```
torch.Size([2, 2])
```

✓ 7.4. Multiple Input and Multiple Output Channels

```
!pip install d2l==1.0.3
```


 숨겨진 출력 표시

```
import torch
from d2l import torch as d2l

def corr2d_multi_in(X, K):
    # Iterate through the 0th dimension (channel) of K first, then add them up
    return sum(d2l.corr2d(x, k) for x, k in zip(X, K))
```


```
X = torch.tensor([[[[0.0, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]],
                    [[1.0, 2.0, 3.0], [4.0, 5.0, 6.0], [7.0, 8.0, 9.0]]]])
K = torch.tensor([[[0.0, 1.0], [2.0, 3.0]], [[1.0, 2.0], [3.0, 4.0]]])
```

```
corr2d_multi_in(X, K)
```


 tensor([[56., 72.],
[104., 120.]])

```
def corr2d_multi_in_out(X, K):
    # Iterate through the 0th dimension of K, and each time, perform
    # cross-correlation operations with input X. All of the results are
    # stacked together
    return torch.stack([corr2d_multi_in(X, k) for k in K], 0)
```

```
K = torch.stack((K, K + 1, K + 2), 0)
K.shape
```

 torch.Size([3, 2, 2, 2])

```
corr2d_multi_in_out(X, K)
```

 tensor([[[[56., 72.],
[104., 120.]],

[[76., 100.],
[148., 172.]],

[[96., 128.],
[192., 224.]]])

```
def corr2d_multi_in_out_1x1(X, K):
    c_i, h, w = X.shape
    c_o = K.shape[0]
    X = X.reshape((c_i, h * w))
    K = K.reshape((c_o, c_i))
    # Matrix multiplication in the fully connected layer
    Y = torch.matmul(K, X)
    return Y.reshape((c_o, h, w))
```

```
X = torch.normal(0, 1, (3, 3, 3))
K = torch.normal(0, 1, (2, 3, 1, 1))
Y1 = corr2d_multi_in_out_1x1(X, K)
Y2 = corr2d_multi_in_out(X, K)
assert float(torch.abs(Y1 - Y2).sum()) < 1e-6
```

✓ 7.5. Pooling

```
!pip install d2l==1.0.3
```

 숨겨진 출력 표시

```
import torch
from torch import nn
from d2l import torch as d2l
```

```
def pool2d(X, pool_size, mode='max'):
    p_h, p_w = pool_size
    Y = torch.zeros((X.shape[0] - p_h + 1, X.shape[1] - p_w + 1))
    for i in range(Y.shape[0]):
        for j in range(Y.shape[1]):
            if mode == 'max':
                Y[i, j] = X[i: i + p_h, j: j + p_w].max()
            elif mode == 'avg':
                Y[i, j] = X[i: i + p_h, j: j + p_w].mean()
    return Y
```

```
X = torch.tensor([[0.0, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]])
pool2d(X, (2, 2))
```

```
tensor([[4., 5.],
        [7., 8.]])
```

```
pool2d(X, (2, 2), 'avg')
```

```
tensor([[2., 3.],
        [5., 6.]])
```

```
X = torch.arange(16, dtype=torch.float32).reshape((1, 1, 4, 4))
X
```

```
tensor([[[[ 0.,  1.,  2.,  3.],
           [ 4.,  5.,  6.,  7.],
           [ 8.,  9., 10., 11.],
           [12., 13., 14., 15.]]]]])
```

```
pool2d = nn.MaxPool2d(3)
# Pooling has no model parameters, hence it needs no initialization
pool2d(X)
```

```
tensor([[[[10.]]]])
```

```
pool2d = nn.MaxPool2d(3, padding=1, stride=2)
pool2d(X)
```

```
tensor([[[[ 5.,  7.],
           [13., 15.]]]]])
```

```
pool2d = nn.MaxPool2d((2, 3), stride=(2, 3), padding=(0, 1))
pool2d(X)
```

```
tensor([[[[ 5.,  7.],
           [13., 15.]]]]])
```

```
X = torch.cat((X, X + 1), 1)
X
```

```
tensor([[[[ 0.,  1.,  2.,  3.],
           [ 4.,  5.,  6.,  7.],
           [ 8.,  9., 10., 11.],
           [12., 13., 14., 15.],

           [ 1.,  2.,  3.,  4.],
           [ 5.,  6.,  7.,  8.],
           [ 9., 10., 11., 12.],
           [13., 14., 15., 16.]]]]])
```

```
pool2d = nn.MaxPool2d(3, padding=1, stride=2)
pool2d(X)
```

```
tensor([[[[ 5.,  7.],
           [13., 15.],

           [ 6.,  8.],
           [14., 16.]]]]])
```

✓ 7.6. Convolutional Neural Networks (LeNet)

```
!pip install d2l==1.0.3
```

 숨겨진 출력 표시

```

import torch
from torch import nn
from d2l import torch as d2l

def init_cnn(module):
    """Initialize weights for CNNs."""
    if type(module) == nn.Linear or type(module) == nn.Conv2d:
        nn.init.xavier_uniform_(module.weight)

class LeNet(d2l.Classifier):
    """The LeNet-5 model."""
    def __init__(self, lr=0.1, num_classes=10):
        super().__init__()
        self.save_hyperparameters()
        self.net = nn.Sequential(
            nn.LazyConv2d(6, kernel_size=5, padding=2), nn.Sigmoid(),
            nn.AvgPool2d(kernel_size=2, stride=2),
            nn.LazyConv2d(16, kernel_size=5), nn.Sigmoid(),
            nn.AvgPool2d(kernel_size=2, stride=2),
            nn.Flatten(),
            nn.LazyLinear(120), nn.Sigmoid(),
            nn.LazyLinear(84), nn.Sigmoid(),
            nn.LazyLinear(num_classes))

@d2l.add_to_class(d2l.Classifier)
def layer_summary(self, X_shape):
    X = torch.randn(*X_shape)
    for layer in self.net:
        X = layer(X)
        print(layer.__class__.__name__, 'output shape:\t', X.shape)

model = LeNet()
model.layer_summary((1, 1, 28, 28))

```

```


↗ Conv2d output shape:      torch.Size([1, 6, 28, 28])
  Sigmoid output shape:     torch.Size([1, 6, 28, 28])
  AvgPool2d output shape:   torch.Size([1, 6, 14, 14])
  Conv2d output shape:      torch.Size([1, 16, 10, 10])
  Sigmoid output shape:     torch.Size([1, 16, 10, 10])
  AvgPool2d output shape:   torch.Size([1, 16, 5, 5])
  Flatten output shape:     torch.Size([1, 400])
  Linear output shape:      torch.Size([1, 120])
  Sigmoid output shape:     torch.Size([1, 120])
  Linear output shape:      torch.Size([1, 84])
  Sigmoid output shape:     torch.Size([1, 84])
  Linear output shape:      torch.Size([1, 10])

```

```

trainer = d2l.Trainer(max_epochs=10, num_gpus=1)
data = d2l.FashionMNIST(batch_size=128)
model = LeNet(lr=0.1)
model.apply_init([next(iter(data.get_dataloader(True)))[0]], init_cnn)
trainer.fit(model, data)

```



✓ 8.2. Networks Using Blocks (VGG)

```
!pip install d2l==1.0.3
```

↗ 숨겨진 출력 표시

```

import torch
from torch import nn

```

```

from d2l import torch as d2l

def vgg_block(num_convs, out_channels):
    layers = []
    for _ in range(num_convs):
        layers.append(nn.LazyConv2d(out_channels, kernel_size=3, padding=1))
        layers.append(nn.ReLU())
    layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
    return nn.Sequential(*layers)

class VGG(d2l.Classifier):
    def __init__(self, arch, lr=0.1, num_classes=10):
        super().__init__()
        self.save_hyperparameters()
        conv_blks = []
        for (num_convs, out_channels) in arch:
            conv_blks.append(vgg_block(num_convs, out_channels))
        self.net = nn.Sequential(
            *conv_blks, nn.Flatten(),
            nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
            nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
            nn.LazyLinear(num_classes))
        self.net.apply(d2l.init_cnn)

```

```

VGG(arch=((1, 64), (1, 128), (2, 256), (2, 512), (2, 512))).layer_summary(
    (1, 1, 224, 224))

```

```


Sequential output shape:      torch.Size([1, 64, 112, 112])
Sequential output shape:      torch.Size([1, 128, 56, 56])
Sequential output shape:      torch.Size([1, 256, 28, 28])
Sequential output shape:      torch.Size([1, 512, 14, 14])
Sequential output shape:      torch.Size([1, 512, 7, 7])
Flatten output shape:         torch.Size([1, 25088])
Linear output shape:          torch.Size([1, 4096])
ReLU output shape:            torch.Size([1, 4096])
Dropout output shape:         torch.Size([1, 4096])
Linear output shape:          torch.Size([1, 4096])
ReLU output shape:            torch.Size([1, 4096])
Dropout output shape:         torch.Size([1, 4096])
Linear output shape:          torch.Size([1, 10])

```

```

model = VGG(arch=((1, 16), (1, 32), (2, 64), (2, 128), (2, 128)), lr=0.01)
trainer = d2l.Trainer(max_epochs=10, num_gpus=1)
data = d2l.FashionMNIST(batch_size=128, resize=(224, 224))
model.apply_init([next(iter(data.get_dataloader(True)))[0]], d2l.init_cnn)
trainer.fit(model, data)

```



✓ 8.6. Residual Networks (ResNet) and ResNeXt

```
!pip install d2l==1.0.3
```

 숨겨진 출력 표시

```

import torch
from torch import nn
from torch.nn import functional as F
from d2l import torch as d2l

```

```

class Residual(nn.Module):
    """The Residual block of ResNet models."""
    def __init__(self, num_channels, use_1x1conv=False, strides=1):

```

```

    super().__init__()
    self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                               stride=strides)
    self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1)
    if use_1x1conv:
        self.conv3 = nn.LazyConv2d(num_channels, kernel_size=1,
                                     stride=strides)
    else:
        self.conv3 = None
    self.bn1 = nn.LazyBatchNorm2d()
    self.bn2 = nn.LazyBatchNorm2d()

    def forward(self, X):
        Y = F.relu(self.bn1(self.conv1(X)))
        Y = self.bn2(self.conv2(Y))
        if self.conv3:
            X = self.conv3(X)
        Y += X
        return F.relu(Y)

```

```

blk = Residual(3)
X = torch.randn(4, 3, 6, 6)
blk(X).shape

```

```

→ torch.Size([4, 3, 6, 6])

```

```

blk = Residual(6, use_1x1conv=True, strides=2)
blk(X).shape

```

```

→ torch.Size([4, 6, 3, 3])

```

```

class ResNet(d2l.Classifier):
    def b1(self):
        return nn.Sequential(
            nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
            nn.LazyBatchNorm2d(), nn.ReLU(),
            nn.MaxPool2d(kernel_size=3, stride=2, padding=1))

@d2l.add_to_class(ResNet)
def block(self, num_residuals, num_channels, first_block=False):
    blk = []
    for i in range(num_residuals):
        if i == 0 and not first_block:
            blk.append(Residual(num_channels, use_1x1conv=True, strides=2))
        else:
            blk.append(Residual(num_channels))
    return nn.Sequential(*blk)

```

```

@d2l.add_to_class(ResNet)
def __init__(self, arch, lr=0.1, num_classes=10):
    super(ResNet, self).__init__()
    self.save_hyperparameters()
    self.net = nn.Sequential(self.b1())
    for i, b in enumerate(arch):
        self.net.add_module(f'b{i+2}', self.block(*b, first_block=(i==0)))
    self.net.add_module('last', nn.Sequential(
        nn.AdaptiveAvgPool2d((1, 1)), nn.Flatten(),
        nn.LazyLinear(num_classes)))
    self.net.apply(d2l.init_cnn)

```

```

class ResNet18(ResNet):
    def __init__(self, lr=0.1, num_classes=10):
        super().__init__((2, 64), (2, 128), (2, 256), (2, 512)),
                        lr, num_classes)

```

```

ResNet18().layer_summary((1, 1, 96, 96))

```

```

→ Sequential output shape:      torch.Size([1, 64, 24, 24])
Sequential output shape:      torch.Size([1, 64, 24, 24])
Sequential output shape:      torch.Size([1, 128, 12, 12])
Sequential output shape:      torch.Size([1, 256, 6, 6])
Sequential output shape:      torch.Size([1, 512, 3, 3])
Sequential output shape:      torch.Size([1, 10])


```

```

model = ResNet18(lr=0.01)
trainer = d2l.Trainer(max_epochs=10, num_gpus=1)
data = d2l.FashionMNIST(batch_size=128, resize=(96, 96))

```

```
model.apply_init([next(iter(data.get_dataloader(True)))[0]], d2l.init_cnn)
trainer.fit(model, data)
```



✓ Discussions & Exercises

✓ 7.3. Padding and Stride

CNN의 input width/height가 주어졌을 때, output width/height를 구하는 공식은 다음과 같다.

O : output width/height

I : input width/height

K : kernel size(width, height)

P : padding size

S : stride

$$O = \frac{I - K + 2P}{S} + 1$$

만약 $I - K + 2P$ 가 S 로 나누어 떨어지지 않는다면 어떻게 될까?

```
import torch
from torch import nn
```

```
# CASE 1
input = torch.zeros(3, 5, 5)
conv2d = nn.Conv2d(3, 3, kernel_size=3, stride=3, padding=1)
# 0 = (5 - 3 + 2) / 3 + 1 = 2.333
```

```
output = conv2d(input)
print(output.shape)
```

```
→ torch.Size([3, 2, 2])
```

```
# CASE 2
input = torch.zeros(3, 7, 7)
conv2d = nn.Conv2d(3, 3, kernel_size=3, stride=4, padding=1)
# 0 = (7 - 3 + 2) / 4 + 1 = 2.5
```

```
output = conv2d(input)
print(output.shape)
```

```
→ torch.Size([3, 2, 2])
```

나누어 떨어지지 않는다면 소수점을 버리게 된다. 가령 case 1의 경우, 주어진 input의 row(또는 column)에 대해 convolution 연산을 2번 시행한 뒤 남은 칸이 발생하는데, 이것은 건너뛰게 되어 최종적으로 width와 height는 2가 되는 것이다.

```
input = torch.tensor([[1, 1, -999999], [1, 1, -999999], [-999999, -999999, -999999.0]]).reshape(1, 3, 3)
conv2d = nn.Conv2d(1, 1, kernel_size=2, stride=2)
print('weight:', conv2d.weight.data)
print('bias:', conv2d.bias.data)
print('output:', conv2d(input))
print('sum of weight and bias:', conv2d.weight.sum() + conv2d.bias.sum())
```

```
→ weight: tensor([[[[-0.4835, -0.1470],
                    [-0.1025, -0.3671]]]])
    bias: tensor([-0.4484])
```



```
output: tensor([[-1.5484]]), grad_fn=<SqueezeBackward1>)
sum of weight and bias: tensor(-1.5484, grad_fn=<AddBackward0>)
```

주어진 3*3 크기의 input의 좌상단 2*2 부분(모든 값이 1)에 대해서 2*2 size의 kernel로 convolution 연산을 시행하고, 남은 부분(모든 값이 -999999)은 건너뛰는 것을 확인할 수 있다.

✓ 7.5. Pooling

Pooling에는 Max pooling, average pooling, min pooling 등이 있다. 이들의 차이점은 무엇이고, 각각을 어떤 상황에 사용해야 할까?

Max pooling: 주변 영역의 maximum pixel 값이 선택된다.

Min pooling: 주변 영역의 minimum pixel 값이 선택된다.



Average pooling: 주변 영역의 pixel 값들의 평균이 선택된다.

이들의 차이점은 다음 그림을 통해 확인할 수 있다.







Average pooling은 이미지를 smooth하여 sharp features가 사라지게 된다.

반면 max pooling은 이미지의 밝은 픽셀을 선택하게 된다. 그렇기에 MNIST dataset과 같이, background가 검은색이고 object가 하얀색인 이미지에서 max pooling이 사용될 수 있다.



이를 통해 computational cost를 줄일 수 있다.

반대로 min pooling은 이미지의 어두운 픽셀을 선택하게 되기에, background가 하얀색이고 object가 검은색인 이미지에 대해 사용될 수 있다.



이처럼 각각의 pooling method는 각자 사용되는 용도가 다르다. 일반적으로 max pooling은 이미지에서 중요한 정보를 추출하기 위해 사용되고, average pooling은 feature의 일반적인 정보를 추출할 때 사용된다. CNN에서는 주로 max pooling이 사용되고, average pooling은 주로 마지막 레이어에서 사용된다. Min pooling은 자주 쓰이지 않는다.